

December 2, 2022

# 1 Exercise Sheet 6

Homework by: Jakob Kreft and Lukas Matt

```
[3]: import numpy as np
from matplotlib import pyplot as plt
import scipy.sparse
from sklearn.linear_model import Ridge, Lasso, LinearRegression
```

## 1.1 1 Regularization and Bias

$$y = \beta_0 + \beta^T x + \epsilon = \tilde{\beta}^T \tilde{x} + \epsilon$$

Consider a regression problem with two explanatory variables  $x_1, x_2$ , i.e.  $\beta = (\beta_0, \beta_1, \beta_2)^T$  and  $x = (1, x_0, x_1)^T$ .

### 1.1.1 (a)

In this setting, write down the loss function for ridge regression, penalizing the L2 -norm of  $\beta$ , in components. What is the influence of the regression strength on the bias  $\beta_0$ ?

$$y = (y_1, y_2, y_3) \in \mathbb{R}^3$$

$$\text{Loss} = \|y - \beta^T x\|^2 + \lambda \|\beta\|^2$$

$$= \|y - \beta^T x\|^2 + \lambda(\beta_0^2 + \beta_1^2 + \beta_2^2)$$

$$= \|y - \beta^T x\|^2 + \lambda(\beta_1^2 + \beta_2^2) + \lambda\beta_0^2$$

The regularization strengt  $\lambda$  is a factor in front of  $\beta_0^2$ . If the regularization is strong, high values of  $\beta_0$  are penalized (the same is true for all  $\beta_i$ ).

### 1.1.2 (b)

Oftentimes, a regularization of the bias term is unwanted. How would you modify the loss function to account for this?

Somehow,  $\beta_0$  needs to be removed / the factor set to  $= 0$ . This can be achieved for example by dealing with  $\beta_0$  separately:

$$y = \beta_0 + \beta^T x + \epsilon$$

where now  $\beta = (\beta_1, \beta_2)^T$  and  $x = (x_0, x_1)^T$ .

The new loss function is:

$$\text{Loss}_{\text{bias}} = \|y - \beta^T x - \beta_0\|^2 + \lambda \|\beta\|^2$$

Alternatively, one might instead introduce the regression strength as a vector quantity  $\vec{\lambda} = (\lambda_i)_i^T$  which allows to set the regularization strength individually for each parameter  $\beta_i$ . The new loss would be:

$$\text{Loss}_{\text{vect}} = \|y - \beta^T x - \beta_0\|^2 + (\vec{\lambda} \circ \beta)^T \beta$$

where  $a \circ b$  is the Hadamard product (element-wise multiplication) of two vectors  $a$  and  $b$ .

Setting  $\lambda_0 = 0$  and  $\lambda_1 = \lambda_2 = \lambda$  in the example at hand yields:

$$\begin{aligned} \text{Loss}_{\text{vect}} &= \|y - \beta^T x\|^2 + \begin{pmatrix} \lambda_0 \beta_0 \\ \lambda_1 \beta_1 \\ \lambda_2 \beta_2 \end{pmatrix}^T \cdot \begin{pmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \end{pmatrix} \\ &= \|y - \beta^T x\|^2 + [0 \cdot \beta_0^2 + \lambda_1 \beta_1^2 + \lambda_2 \beta_2^2] \end{aligned}$$

which results in the same loss function as above.

### 1.1.3 (c)

Which shapes in  $\mathbb{R}^3$  do the regularization contours (i.e. sets of parameters with equal regularization penalty) of versions (a) and (b) have?

For ridge regression, the regularization term is an added squared L2-norm of  $\beta$ . We consider the space spanned by  $\beta = (\beta_0, \beta_1, \beta_2)^T \in \mathbb{R}^3$ .

In (a), the regularization contours are just the surfaces for which the squared L2-norm gives equal values:

$$\{\beta \in \mathbb{R}^3 : \|\beta\|_2^2 = g\}, \quad \text{with } g \text{ any fixed value} \in [0, \infty)$$

which form a (hyper-) sphere.

For (b), the first dimension collapses since  $\lambda_0 = 0$  and the resulting set of points:

$$\{\beta \in \mathbb{R}^3 : \beta_1^2 + \beta_2^2 = g\}$$

takes the shape of a circle in the y-z-plane.

## 1.2 2 Estimating Parameter Relevance

```
[4]: # load the data
with open('data/vostok.txt', 'r') as f:
    lines = f.readlines()

# remove header and split lines
lines = [l.split() for l in lines[2:]]

# filter out lines with missing data
lines = [l for l in lines if len(l) == 4]

# convert to float
lines = np.array(lines).astype(np.float32)
print(f'{lines.shape=}')

features = np.concatenate([lines[:, :1], lines[:, 2:]], axis=1).T
feature_names = 'age', 'CO ', 'dust'
labels = lines[:, 1]
label_name = 'T'

print(f'{features.shape=}, {labels.shape=}')
```

```
lines.shape=(3729, 4)
features.shape=(3, 3729), labels.shape=(3729,)
```

```
[5]: from sklearn.linear_model import LinearRegression

# TODO: fit the linear regressor and compute the sum of square deviations

def get_linreg_ssd(X, Y):

    regressor = LinearRegression()
    regressor.fit(X.T, Y)

    pred_Y = regressor.predict(X.T)

    ssd = np.sum((pred_Y - Y)**2)
    return ssd
```

```
[6]: original = get_linreg_ssd(features, labels)
print('sum of squared deviations:', original)
```

sum of squared deviations: 6362.9375

```
[7]: # TODO: for each feature, randomly permute it amongst the samples,
#      refit the regressor and compute sum of squared deviations
```

```
rng = np.random.default_rng()
ssds = []

for p, row in enumerate(features):

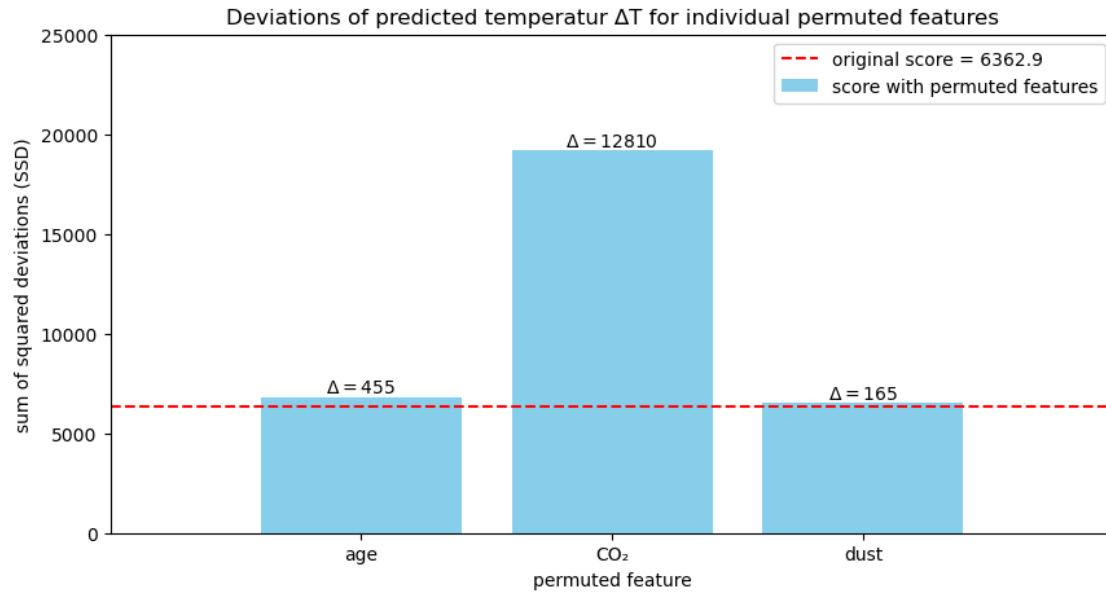
    perm_features = features.copy()

    perm_features[p] = rng.permutation(row)

    ssds.append(get_linreg_ssd(perm_features, labels))
```

```
[8]: plt.figure(figsize=(10,5))
plt.title('Deviations of predicted temperatur T for individual permuted_
↳features')
bar = plt.bar([0,1,2], ssds, color='skyblue', label='score with permuted_
↳features')
plt.plot([-1,3], np.ones(2)*original, 'r--', label=r'original score =_
↳'+str(round(original, 1)))
for rect in bar:
    delta = rect.get_height() - original
    plt.text(rect.get_x() + rect.get_width() / 2.0, rect.get_height(),_
↳f'$\Delta={delta:.0f}$', ha='center', va='bottom')

plt.xticks([0,1,2], feature_names)
plt.xlabel('permuted feature')
plt.xlim(-1, 3)
plt.ylabel('sum of squared deviations (SSD)')
plt.ylim(0, 25000)
plt.legend()
plt.show()
```



The resulting sum of squared deviation scores are plotted as a bar chart. The red line marks the original ssd score.

The deviation is strongest for the permuted CO<sub>2</sub> feature, which suggests that it has the highest relevance for predicting the temperature  $T$ .

The score is slightly increased for the age feature, and there is almost no difference for the dust feature, which therefore appears to be the least relevant.

### 1.3 4 Visualize Regularization Contours

```
[9]: # load the data
data = np.load('data/linreg.npz')
x = data['X']
y = data['Y']
print(f'{x.shape} {y.shape}')
```

```
(2, 100) (1, 100)
```

```
[10]: # TODO: create a grid of points in the parameter space
beta_i = np.linspace(-1, 3, 500)
betagrid = np.array(np.meshgrid(beta_i, beta_i))
```

#### 1.3.1 (a)

Plot the Ridge regression regularization term as well as the Lasso regularization term for  $\beta_1, \beta_2 \in [-1, 3]$ .

Ridge regression uses  $L_2$  as regularization term:  $\|\beta\|_2^2$

Lasso regression used  $L_1$  as regularization term:  $\|\beta\|_1$

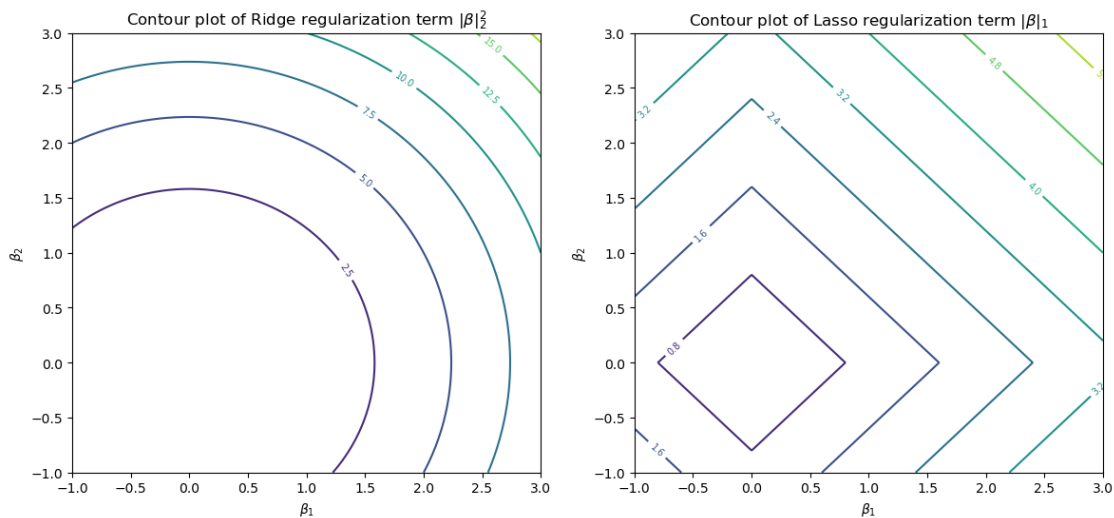
```
[11]: def ridge_regterm(b1, b2):
        return b1**2 + b2**2

def lasso_regterm(b1, b2):
    return np.abs(b1) + np.abs(b2)

[12]: # TODO: make contour plots for ridge and lasso regularization terms
fig, ax = plt.subplots(1,2, figsize=(14,6))
beta0, beta1 = betagrid

ax[0].set_title(r'Contour plot of Ridge regularization term  $\|\beta\|_2^2$ ')
ridge = ax[0].contour(beta0, beta1, ridge_regterm(beta0, beta1))
ax[0].clabel(ridge, inline=True, fontsize=7)
ax[0].set_xlabel(r'$\beta_1$')
ax[0].set_ylabel(r'$\beta_2$')

ax[1].set_title(r'Contour plot of Lasso regularization term  $\|\beta\|_1$ ')
lasso = ax[1].contour(beta0, beta1, lasso_regterm(beta0, beta1))
ax[1].clabel(lasso, inline=True, fontsize=7)
ax[1].set_xlabel(r'$\beta_1$')
ax[1].set_ylabel(r'$\beta_2$')
plt.show()
```



### 1.3.2 (b)

For the data set linreg.npz plot the sum of squares (SSQ) of a linear regression as a function of  $\beta$  over the same range as in (a), i.e. over the grid  $[-1, 3] \times [-1, 3]$ .

```
[13]: print(x.shape)
      print(y.shape)
```

```
(2, 100)
(1, 100)
```

```
[17]: # TODO: for each combination of parameters, compute the sum of squared
      ↪ deviations.
      # do not use loops, but numpy broadcasting!
      # TODO: make a contour plot for sum of squared deviations

      # intuitive solution using one loop:
      lypred = np.array([beta0 * x[0][n] + beta1 * x[1][n] for n in range(x.
      ↪ shape[1])])
      lssq = np.sum((lypred - np.expand_dims(y.T, axis=1))**2, axis=0)
      # we used this to verify the implementation below
```

```
[18]: # We use the dot product to compute  $\beta^T X$ 
      # To match shapes, we need to transpose. Remember this for later
      print(f'{betagrid.shape = }')
      ypred = np.dot(np.array(betagrid).T, x)
      print(ypred.shape)
```

```
betagrid.shape = (2, 500, 500)
(500, 500, 100)
```

```
[19]: # just to visualize that broadcasting works out here
      (ypred - y).shape
```

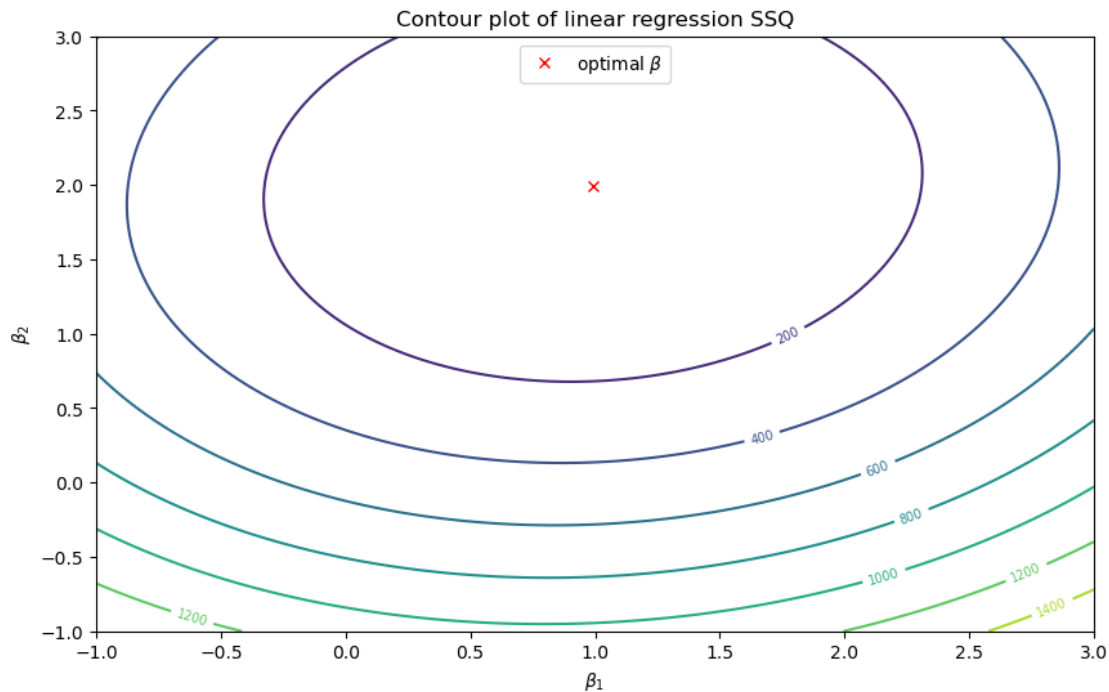
```
[19]: (500, 500, 100)
```

```
[20]: # calculate sum of squared errors
      ssq = np.sum((ypred - y)**2, axis=2).T
```

```
[21]: # We run the linear regression and return the optimal parameters to see if our
      ↪ plot makes sense
      beta = LinearRegression().fit(x.T, y.T).coef_[0]
```

```
[99]: plt.figure(figsize=(10,6))
      ax = plt.gca()
      ax.set_title(r'Contour plot of linear regression SSQ')
      ridge = ax.contour(beta0, beta1, ssq)
      ax.clabel(ridge, inline=True, fontsize=7)
      plt.plot(*beta, 'r x', label=r'optimal  $\beta$ ')
      ax.set_xlabel(r' $\beta_1$ ')
      ax.set_ylabel(r' $\beta_2$ ')
      plt.legend(loc='upper center')
```

```
plt.show()
```



```
[23]: for i, b in enumerate(beta):  
       print('beta'+str(i), b)
```

```
beta0 0.9912379503154678
```

```
beta1 1.9922619503237609
```

We mark the position of the optimal parameter set  $\beta$  obtained by linear regression (red x). It agrees well with the plotted contours

### 1.3.3 (c)

Plot the ridge and Lasso loss functions, i.e.  $\text{SSQ}(\beta) + \lambda\|\beta\|_2^2$  and  $\text{SSQ}(\beta) + \lambda\|\beta\|_1$ , for  $\lambda \in 0, 10, 50, 100, 200, 300$  in the same  $\beta$  grid as before and discuss your observations!

```
[24]: # TODO: for each lambda, plot both ridge regression and lasso loss functions  
       lambdas = [0, 10, 50, 100, 200, 300]
```

## 1.4 5 CT Reconstruction

set up design matrix (run this once to save to disk)

```
[ ]: # create design matrix  
     # don't change any of this, just run it once to create and save the design  
     ↪matrix
```



```

import os

if not os.path.exists('data/design_matrix.npy'):
    res = (99, 117)
    xs = np.arange(0, res[1]+1) - res[1]/2 # np.linspace(-1, 1, res[1] + 1)
    ys = np.arange(0, res[0]+1) - res[0]/2 # np.linspace(-1, 1, res[0] + 1)

    # rays are defined by origin and direction
    n_parallel_rays = 70
    ray_offset_range = [-res[1]/1.5, res[1]/1.5]
    n_ray_angles = 30
    n_rays = n_parallel_rays * n_ray_angles

    ray_angles = np.linspace(0, np.pi, n_ray_angles, endpoint=False) + np.pi/
    ↪ n_ray_angles

    # offsets for ray_angle = 0, i.e. parallel to x-axis
    ray_0_offsets = np.stack([np.zeros(n_parallel_rays), np.
    ↪ linspace(*ray_offset_range, n_parallel_rays)], axis=-1)
    ray_0_directions = np.stack([np.ones(n_parallel_rays), np.
    ↪ zeros(n_parallel_rays)], axis=-1)

    def rot_mat(angle):
        c, s = np.cos(angle), np.sin(angle)
        return np.stack([np.stack([c, s], axis=-1), np.stack([-s, c],
    ↪ axis=-1)], axis=-1)

    ray_rot_mats = rot_mat(ray_angles)

    ray_offsets = np.einsum('oi,aij->aoj', ray_0_offsets, ray_rot_mats).
    ↪ reshape(-1, 2)
    ray_directions = np.einsum('oi,aij->aoj', ray_0_directions, ray_rot_mats).
    ↪ reshape(-1, 2)

    sigma = 1
    kernel = lambda x: np.exp(-x**2/sigma**2/2)

    xsc = (xs[1:] + xs[:-1]) / 2
    ysc = (ys[1:] + ys[:-1]) / 2
    b = np.stack(np.meshgrid(xsc, ysc), axis=-1).reshape(-1, 2)
    a = ray_offsets
    v = ray_directions
    v = v / np.linalg.norm(v, axis=-1, keepdims=True)
    p = ((b[None] - a[:, None]) * v[:, None]).sum(-1, keepdims=True) * v[:,
    ↪ None] + a[:, None]
    d = np.linalg.norm(b - p, axis=-1)

```

```

d = kernel(d)
design_matrix = d.T

np.save('data/design_matrix.npy', design_matrix)
print(f'created and saved design matrix of shape {design_matrix.shape} at_
↳data/design_matrix.npy')

```

One application of linear regression is the reconstruction of CT-Scans. In this task, you will do this on simulated data in the 2D case. You are given a sinogram  $Y \in \mathbb{R}^{ar}$ , a matrix where each row corresponds to a (1D) projection of the image consisting of  $r$  detector readouts along one of a distinct, evenly spaced angles. Additionally, you are given the design matrix  $X \in \mathbb{R}^{p \times ar}$ . Excluding noise, one has  $Y = IX$ , with the image  $I \in \mathbb{R}^p$  which should be reconstructed.

Some facts/descriptions for our understanding:

Each 1D projection consists  $r$  detector pixel readouts, and there are  $a$  different, evenly-spaced projection angles.

Rows of the sinogram  $Y \in \mathbb{R}^{ar}$  are projections. Columns are the readout value of one pixel over all the angles.

#### 1.4.1 (a)

What is the interpretation of a column of  $X$ ? Visualize a choice of four columns as images.

```

[67]: design_matrix = np.load('data/design_matrix.npy')
      sino = np.load('data/sino.npy')

      print(f'{design_matrix.shape=}')
      print(f'{sino.shape=}')

```

```

design_matrix.shape=(11583, 2100)
sino.shape=(1, 2100)

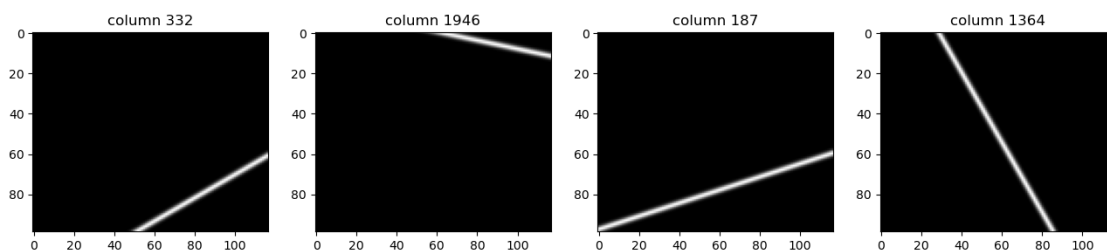
```

```

[68]: # TODO: visualize four random columns as images, using an image shape of (99,
      ↳117)
      img_shape = (99, 117)

      fig, axs = plt.subplots(1, 4, figsize=(16, 4))
      for i, ax in zip(np.random.choice(np.arange(design_matrix.shape[1]), 4), axs):
          ax.imshow(design_matrix[:, i].reshape(*img_shape), cmap='gray');
          ax.set_title(f'column {i}')

```



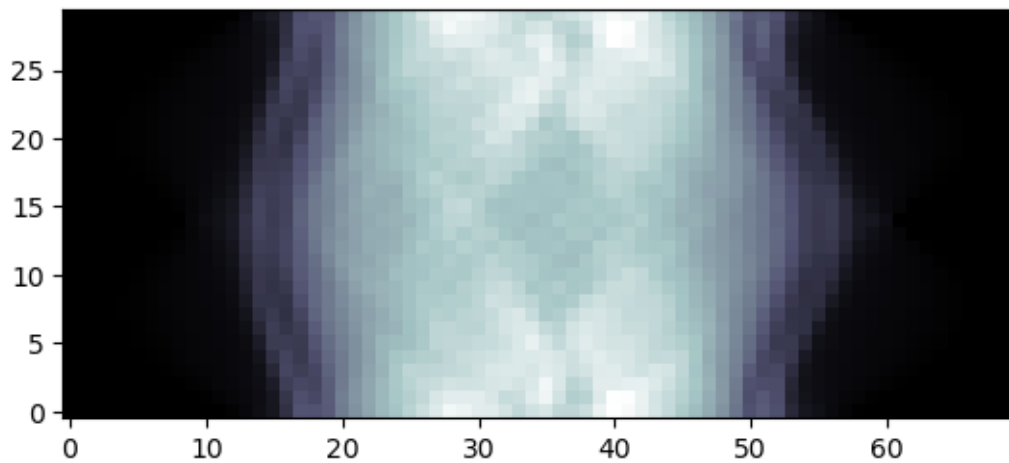
For the design matrix  $X \in \mathbb{R}^{p \times ar}$ :

The rows are the  $99 \times 117 = 11583$  pixels of the resulting image.

The **columns** correspond to the individual x-rays passing through the object at different angles and different detector positions: 70 parallel detector readouts  $\times$  30 angles = 2100 rays.

The matrix describes how much each x-ray intersects with each of the pixels.

```
[92]: # visualize sinogram as image
n_parallel_rays = 70
n_angles = 30
plt.imshow(sino.reshape(n_angles, n_parallel_rays), origin='lower', cmap='bone')
plt.show();
```



### 1.4.2 (b)

Solve the reconstruction problem with linear regression without any regularization and with ridge regression. What do you observe?

The reconstruction problem is given by:

$$Y = \beta^T X$$

Where  $Y$  is the sinogram,  $\beta$  is the wanted array of pixels (the image) and  $X$  is the design matrix.

We obtain  $\beta$  by performing from `sklearn.linear_model.LinearRegression` and returning the parameters/coefficients.

```
[86]: # TODO: solve the reconstruction with linear regression and visualize the result
linreg = LinearRegression(copy_X=True)
linreg.fit(sino.T, design_matrix.T)
```

```
# TODO: solve the reconstruction with ridge regression and visualize the result
ridreg = Ridge(copy_X=True)
ridreg.fit(sino.T, design_matrix.T)
```

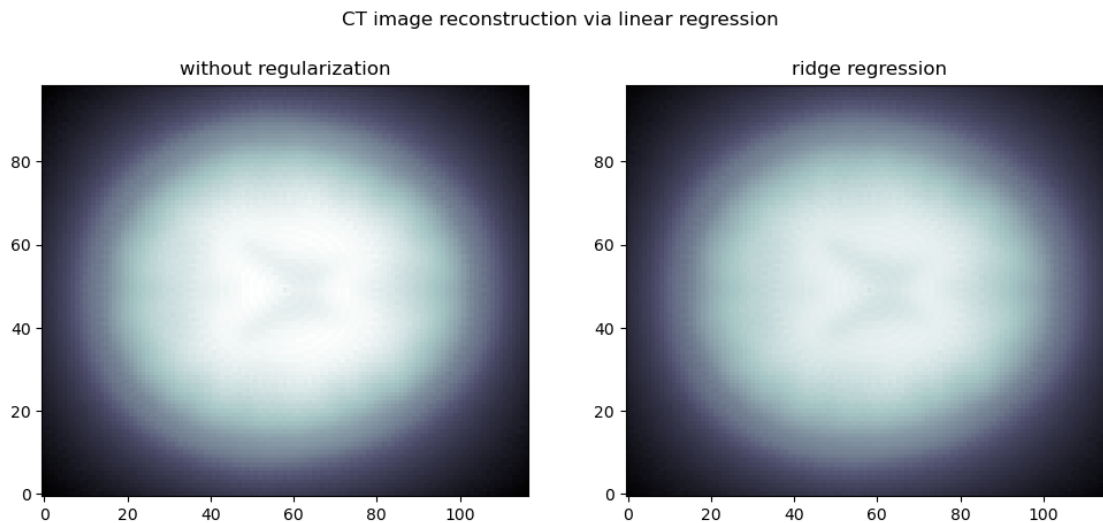
```
linimg = linreg.coef_.reshape(img_shape)
ridimg = ridreg.coef_.reshape(img_shape)
print(f'{linimg.shape = }\n{ridimg.shape = }')
```

```
linimg.shape = (99, 117)
```

```
ridimg.shape = (99, 117)
```

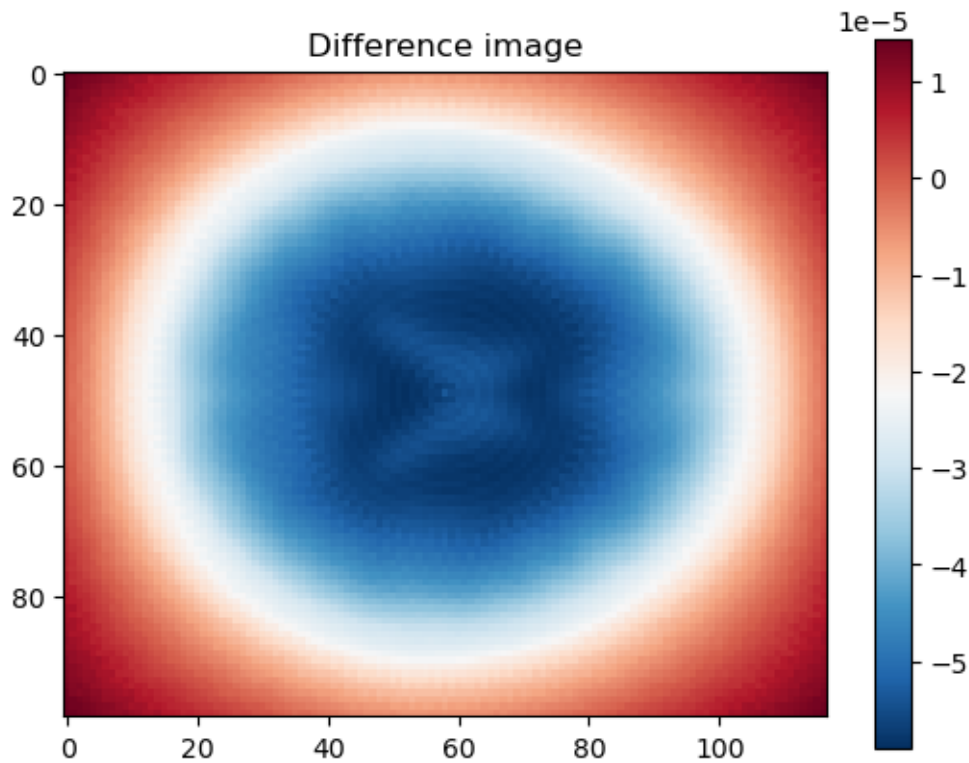
```
[93]: plt.subplots(1,2, figsize=(12,5))
plt.suptitle('CT image reconstruction via linear regression')

vmin, vmax = min(linimg.min(), ridimg.min()), max(linimg.max(), ridimg.max())
plt.subplot(1,2,1)
plt.title('without regularization')
plt.imshow(linimg, cmap='bone', interpolation='none', origin='lower',
           ↪vmin=vmin, vmax=vmax)
plt.subplot(1,2,2)
plt.title('ridge regression')
plt.imshow(ridimg, cmap='bone', interpolation='none', origin='lower',
           ↪vmin=vmin, vmax=vmax)
plt.show()
```



```
[94]: plt.title('Difference image')
plt.imshow(ridimg - linimg, cmap='RdBu_r')
plt.colorbar()
```

```
plt.show()
```



The image reconstructed with ridge regression (regularization) appears overall clearer. It is slightly more sharp/contrasted. The brain matter is more homogenous, while it shows “over-exposed” towards the center in the un-regularized image.

```
[97]: # Optional: try out different regularization strengths and observe the influence
alphas = [0, 1E-1, 1E+1, 1E+2, 1E+3, 1E+4, 1E+5, 1E+6, 1E+7]

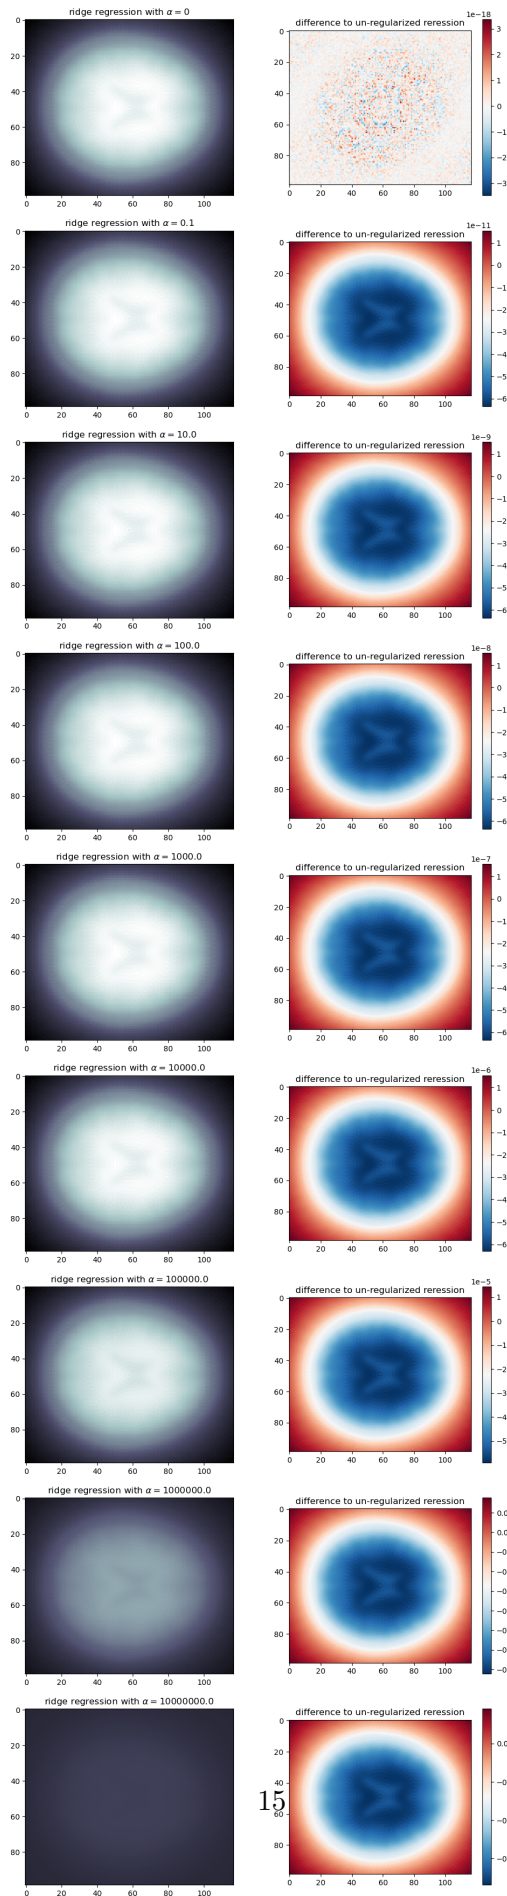
plt.subplots(len(alphas), 2, figsize=(12, 5 * len(alphas)))

for i, a in enumerate(alphas):
    ridreg = Ridge(copy_X=True, alpha=a)
    ridreg.fit(sino.T, design_matrix.T)
    ridimg = ridreg.coef_.reshape(img_shape)
    diffimg = ridimg - lining

    vmin, vmax = min(lining.min(), ridimg.min()), max(lining.max(), ridimg.
    ↪max())
    plt.subplot(len(alphas), 2, i*2+1)
    plt.title(r'ridge regression with $\alpha = ' + str(round(a, 1)) + '$')
    plt.imshow(ridimg, cmap='bone', vmin=vmin, vmax=vmax)
```

```
plt.subplot(len(alphas), 2, i*2+2)
plt.title('difference to un-regularized reression')
plt.imshow(diffimg, cmap='RdBu_r')#, vmin=vmin*1E-5, vmax=vmax*1E-5)
plt.colorbar()

plt.show()
```



[ ]: